

In [1]:

```
import pandas as pd
df=pd.read_csv("Churn_Modelling.csv") # import dataset
print(df)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	
...	...	...	...	...	...	...	...	
9995	9996	15606229	Obijiaku	771	France	Male	39	
9996	9997	15569892	Johnstone	516	France	Male	35	
9997	9998	15584532	Liu	709	France	Female	36	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	
9999	10000	15628319	Walker	792	France	Female	28	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	
...	...	...	...	...	...	
9995	5	0.00	2	1	0	
9996	10	57369.61	1	1	1	
9997	7	0.00	1	0	1	
9998	3	75075.31	2	1	0	
9999	4	130142.79	1	1	0	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0
...	...	...
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0

[10000 rows x 14 columns]

In [30]:


```
#@title 1. Univarient Analysis
```

In [32]:

```
import matplotlib.pyplot as plt
plt.figure(figsize=(5,5))
plt.pie(df['Geography'].value_counts(),startangle=90,autopct='%.3f',labels=['france','spain','germany'])
```

Out[32]:

```
([<matplotlib.patches.Wedge at 0x7faff9505490>,
 <matplotlib.patches.Wedge at 0x7faff9505bd0>,
 <matplotlib.patches.Wedge at 0x7faff950f490>],
 [Text(-1.0999893606763749, -0.004838015996287074, 'france'),
 Text(0.786805947043686, -0.7687238787085312, 'spain'),
 Text(0.7721769705773018, 0.7834173384027577, 'germany')],
 [Text(-0.599994196732568, -0.002638917816156585, '50.140'),
 Text(0.4291668802056469, -0.419303933841017, '25.090'),
 Text(0.42118743849671003, 0.427318548219686, '24.770')])
```



In [ ]:

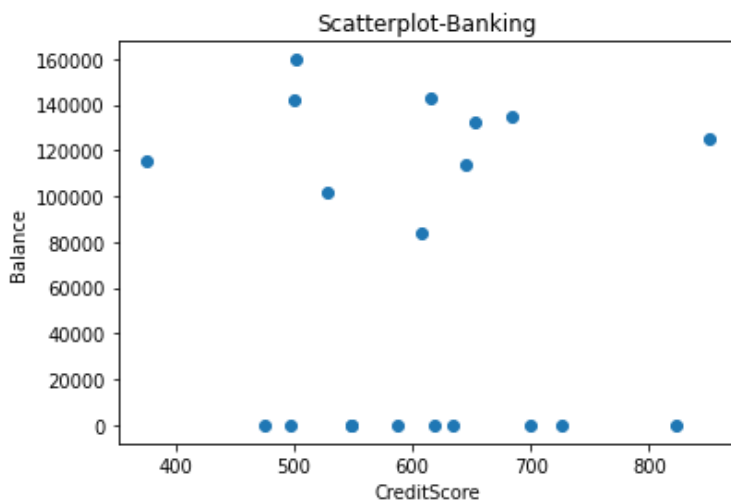
```
#@title 2. Bi - Variate Analysis
```

In [37]:

```
dfs1 = df.head(20)
plt.scatter(dfs1.CreditScore,dfs1.Balance)
plt.title('Scatterplot-Banking')
plt.xlabel("CreditScore")
plt.ylabel("Balance")
```

Out[37]:

Text(0, 0.5, 'Balance')



In [10]:

```
import numpy as np
import seaborn as sns
```

In [28]:

```
#@title 3. Multi - Variate Analysis
```

In [34]:


```
df=sns.catplot(x="Geography",y="EstimatedSalary",hue="Gender",kind="swarm",data=df)
print(df)
```

/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 80.8% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.

warnings.warn(msg, UserWarning)  
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 62.1% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.

```
pplot.  
warnings.warn(msg, UserWarning)  
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 62.5% of  
the points cannot be placed; you may want to decrease the size of the markers or use stri  
pplot.  
warnings.warn(msg, UserWarning)
```

```
<seaborn.axisgrid.FacetGrid object at 0x7faff96fb610>
```



```
In [ ]:
```

```
#@title 4. Perform descriptive statistics on the dataset.
```

```
In [3]:
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 10000 entries, 0 to 9999  
Data columns (total 14 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   RowNumber              10000 non-null  int64  
1   CustomerId             10000 non-null  int64  
2   Surname                10000 non-null  object  
3   CreditScore            10000 non-null  int64  
4   Geography              10000 non-null  object  
5   Gender                 10000 non-null  object  
6   Age                   10000 non-null  int64  
7   Tenure                 10000 non-null  int64  
8   Balance                10000 non-null  float64  
9   NumOfProducts          10000 non-null  int64  
10  HasCrCard              10000 non-null  int64  
11  IsActiveMember         10000 non-null  int64  
12  EstimatedSalary        10000 non-null  float64  
13  Exited                 10000 non-null  int64  
dtypes: float64(2), int64(9), object(3)  
memory usage: 1.1+ MB
```

```
In [4]:
```

```
df.describe()
```

```
Out[4]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	Is
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	

std	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	Exited
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.000000	1.000000	0.000000			
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.000000	1.000000	0.000000			
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.000000	1.000000				
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.000000	1.000000				
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.000000	1.000000				

In [9]:

```
#@title 5. Handle the Missing values.
```

In [7]:

```
values={"RowNumber":0, "CustomerId":0, "CreditScore":0, "Age":0, "Tenure":0, "Balance":0
, "NumOfProducts":0, "HasCrCard":0, "IsActiveMember":0, "EstimatedSalary":0, "Exited":0}
df.fillna(value=values)
```

Out[7]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	
...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00	2	
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	

10000 rows x 14 columns

In [ ]:

```
#@title 6. Find the outliers and replace the outliers
```

In [17]:


```
df= pd.read_csv("Churn_Modelling.csv")
sns.boxplot(df.CreditScore)
```


/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[17]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f518dc0c810>





In [18]:


```
Q1= df.CreditScore.quantile(0.25)
Q3=df.CreditScore.quantile(0.75)
IQR=Q3-Q1
upper_limit =Q3 + 1.5*IQR
lower_limit =Q1 - 1.5*IQR
df['CreditScore'] = np.where(df['CreditScore']>upper_limit,30,df['CreditScore'])
sns.boxplot(df.CreditScore)
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[18]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f518db93f10>



In [ ]:

#@title 7. Check for Categorical columns and perform encoding.

In [16]:

```
df.head(5)
```

Out[16]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	

In [22]:

#label encoding

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
df.Gender= le.fit_transform(df.Gender)
df.head(20)
```

Out[22]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrC
0	1	15634602	Hargrave	619	France	0	42	2	0.00	1	
1	2	15647311	Hill	608	Spain	0	41	1	83807.86	1	
2	3	15619304	Onio	502	France	0	42	8	159660.80	3	
3	4	15701354	Boni	699	France	0	39	1	0.00	2	
4	5	15737888	Mitchell	850	Spain	0	43	2	125510.82	1	
5	6	15574012	Chu	645	Spain	1	44	8	113755.78	2	
6	7	15592531	Bartlett	822	France	1	50	7	0.00	2	
7	8	15656148	Obinna	376	Germany	0	29	4	115046.74	4	
8	9	15792365	He	501	France	1	44	4	142051.07	2	
9	10	15592389	H?	684	France	1	27	2	134603.88	1	
10	11	15767821	Bearce	528	France	1	31	6	102016.72	2	
11	12	15737173	Andrews	497	Spain	1	24	3	0.00	2	
12	13	15632264	Kay	476	France	0	34	10	0.00	2	
13	14	15691483	Chin	549	France	0	25	5	0.00	2	
14	15	15600882	Scott	635	Spain	0	35	7	0.00	2	
15	16	15643966	Goforth	616	Germany	1	45	3	143129.41	2	
16	17	15737452	Romeo	653	Germany	1	58	1	132602.88	1	
17	18	15788218	Henderson	549	Spain	0	24	9	0.00	2	
18	19	15661507	Muldrow	587	Spain	1	45	6	0.00	1	
19	20	15568982	Hao	726	France	0	24	6	0.00	2	

In [24]:

```
#one hot encoding
df_main=pd.get_dummies(df,columns=['Geography'])
df_main.head()
```

Out[24]:

	RowNumber	CustomerId	Surname	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveM
0	1	15634602	Hargrave	619	0	42	2	0.00	1	1	
1	2	15647311	Hill	608	0	41	1	83807.86	1	0	
2	3	15619304	Onio	502	0	42	8	159660.80	3	1	
3	4	15701354	Boni	699	0	39	1	0.00	2	0	
4	5	15737888	Mitchell	850	0	43	2	125510.82	1	1	

In [ ]:

```
#@title 8. Split the data into dependent and independent variables.
```

In [25]:

```
#Splitting the Dataset into the Independent Feature Matrix
df=pd.read_csv("Churn_Modelling.csv")
```

```
X = df.iloc[:, :-1].values
print(X)
```

```
[[1 15634602 'Hargrave' ... 1 1 101348.88]
 [2 15647311 'Hill' ... 0 1 112542.58]
 [3 15619304 'Onio' ... 1 0 113931.57]
 ...
 [9998 15584532 'Liu' ... 0 1 42085.58]
 [9999 15682355 'Sabbatini' ... 1 0 92888.52]
 [10000 15628319 'Walker' ... 1 0 38190.78]]
```

In [26]:

```
#Extracting the Dataset to Get the Dependent Vector
Y = df.iloc[:, -1].values
print(Y)
```

```
[1 0 1 ... 1 1 0]
```

In [ ]:

```
#@title 9. Scale the independent variables
```

In [27]:

```
w = df.head()
q = w[['Age', 'Balance', 'EstimatedSalary']] #spliting the dataset into measureable values
q
```

Out[27]:

	Age	Balance	EstimatedSalary
0	42	0.00	101348.88
1	41	83807.86	112542.58
2	42	159660.80	113931.57
3	39	0.00	93826.63
4	43	125510.82	79084.10

In [28]:

```
from sklearn.preprocessing import scale # library for scalling
from sklearn.preprocessing import MinMaxScaler
mm = MinMaxScaler()

x_scaled = mm.fit_transform(q)
x_scaled
```

Out[28]:

```
array([[0.75      , 0.        , 0.63892099],
       [0.5       , 0.52491194, 0.96014087],
       [0.75      , 1.        , 1.        ],
       [0.        , 0.        , 0.42305883],
       [1.        , 0.78610918, 0.        ]])
```

In [30]:

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x_ss = sc.fit_transform(q)
x_ss
```

Out[30]:

```
array([[ 0.44232587, -1.13763618,  0.09337626],
       [-0.29488391,  0.15434425,  0.96285595],
       [ 0.44232587,  1.32369179,  1.07074687],
       [-1.76930347, -1.13763618, -0.49092058],
       [ 1.17953565,  0.79723632, -1.6360585 ]])
```

In [32]:

```
from sklearn.preprocessing import scale
X_scaled=pd.DataFrame(scale(q),columns=q.columns)
X_scale=X_scaled.head()
X_scale
```

Out[32]:

	Age	Balance	EstimatedSalary
0	0.442326	-1.137636	0.093376
1	-0.294884	0.154344	0.962856
2	0.442326	1.323692	1.070747
3	-1.769303	-1.137636	-0.490921
4	1.179536	0.797236	-1.636059

In [ ]:

```
#@title 10. Split the data into training and testing
```

In [33]:

```
x= df[['Age','Balance','EstimatedSalary']]
x
```

Out[33]:

	Age	Balance	EstimatedSalary
0	42	0.00	101348.88
1	41	83807.86	112542.58
2	42	159660.80	113931.57
3	39	0.00	93826.63
4	43	125510.82	79084.10
...	...	...	...
9995	39	0.00	96270.64
9996	35	57369.61	101699.77
9997	36	0.00	42085.58
9998	42	75075.31	92888.52
9999	28	130142.79	38190.78

10000 rows x 3 columns

In [34]:

```
y = df['Balance']
y
```

Out[34]:

0	0.00
1	83807.86
2	159660.80
3	0.00
4	125510.82
...	...
9995	0.00
9996	57369.61
9997	0.00
9998	75075.31
9999	130142.79



Name: Balance, Length: 10000, dtype: float64

In [35]:

```
#scaling
from sklearn.preprocessing import StandardScaler, MinMaxScaler
sc = StandardScaler()
x_scaled1 = sc.fit_transform(x)
x_scaled1
```

Out[35]:

```
array([[ 0.29351742, -1.22584767,  0.02188649],
       [ 0.19816383,  0.11735002,  0.21653375],
       [ 0.29351742,  1.33305335,  0.2406869 ],
       ...,
       [-0.27860412, -1.22584767, -1.00864308],
       [ 0.29351742, -0.02260751, -0.12523071],
       [-1.04143285,  0.85996499, -1.07636976]])
```

In [36]:

```
#train and test data
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x_scaled1, y, test_size = 0.3, random_state = 0)
```

In [37]:

```
x_train
```

Out[37]:

```
array([[ -0.56466489,  1.11721307, -0.77021814],
       [  0.00745665, -1.22584767, -1.39576675],
       [  3.53553951,  1.35419118, -1.49965629],
       ...,
       [-0.37395771,  1.35890908,  1.41441489],
       [-0.08789694, -1.22584767,  0.84614739],
       [  0.86563897,  0.50630343,  0.32630495]])
```

In [38]:

```
x_train.shape
```

Out[38]:

```
(7000, 3)
```

In [39]:

```
x_test
```

Out[39]:

```
array([[ -0.37395771,  0.87532296,  1.61304597],
       [  0.10281024,  0.42442221,  0.49753166],
       [  0.29351742,  0.30292727, -0.4235611 ],
       ...,
       [  0.10281024,  1.46672809,  1.17045451],
       [  2.86806437,  1.25761599, -0.50846777],
       [  0.96099256,  0.19777742, -1.15342685]])
```

In [40]:

```
x_test.shape
```

Out[40]:

```
(3000, 3)
```

In [41]:

```
y_train
```

```
Out[41]:
```

```
7681    146193.60
9031         0.00
3691    160979.68
202         0.00
5625    143262.04
...
9225    120074.97
4859    114440.24
3264    161274.05
9845         0.00
2732    108076.33
Name: Balance, Length: 7000, dtype: float64
```

```
In [42]:
```

```
y_test
```

```
Out[42]:
```

```
9394    131101.04
898     102967.41
2398     95386.82
5906    112079.58
2343    163034.82
...
4004         0.00
7375     80926.02
9307    168001.34
8394    154953.94
5233     88826.07
Name: Balance, Length: 3000, dtype: float64
```

```
In [ ]:
```